

RAG at 10 million documents

What a retrieval-augmented generation system needs once the corpus stops fitting on a laptop, component by component.

~~1,000 docs~~ **10,000,000** documents, messy, duplicated, permissioned

Retrieval-augmented generation replaces trust in a model's memory with a connection to your own data. The textbook version is small: take a folder of about a thousand documents, split them into pieces, turn each piece into an embedding (a numeric fingerprint of its meaning), and store those in a vector database. At question time you embed the query, pull the closest pieces, and hand them to the LLM as context. With LangChain or LlamaIndex that is roughly twenty lines of code and one afternoon, and it runs fine on a laptop.


That design breaks when the corpus becomes ten million real documents. Scanned contracts, nested spreadsheets, Slack exports and screenshots of tables show up. Duplicates multiply. Security rules decide who may see what. Latency budgets get tight, and some users actively try to break the system. The simple pipeline turns into something closer to an air traffic control tower, and every box in it exists to fix one failure mode that only appears at scale.

Contents

Seven sections, in the order data flows through the system — ingestion to evaluation.

The full map	3
1 Ingestion	4
Parsing: formats are the first enemy	4
Chunking: structure-aware cutting	7
Choosing a chunking strategy	9
Metadata: non-negotiable at scale	10
2 Database layer and retrieval	12
Vector stores are approximate; HNSW	12
Embeddings, cosine and the ANN family	13
Hybrid search: vectors + BM25	14
Filtering, re-rank and the funnel	15
Classical primitives, applied to AI workloads	18
3 The reasoning engine	19
Router, multi-agent, feedback loop	19
4 Human validation as risk tiering	23
5 Red teaming the agent brain	24
6 Evaluation	25
7 RAG vs fine-tuning, offline vs online, capacity	26
The three pillars (summary)	27

RAG IN THE REAL WORLD




The full map. Data sources feed ingestion, ingestion feeds the database layer, and the retrieval funnel serves the reasoning engine and its agents. Human validation gates risky actions, red teaming attacks the reasoning layer, and evaluation closes the loop.

The rest of this document walks the map in the order data flows through it: ingestion, the database and retrieval layer, the reasoning engine, human validation, red teaming, and evaluation.

1 Ingestion


Ingestion has three jobs: restructure raw data so it has a consistent shape, chunk it along lines of meaning, and enrich each chunk with metadata. Garbage that gets past this stage cannot be fixed downstream.



Ingestion in three stages. Re-structuring (document parser, structure analyzer), structure-aware chunking (table preserver, heading detector, boundary detector), and metadata creation (summary, keyword and question generators).


Parsing: formats are the first enemy

The first obstacle is not AI or retrieval. It is file formats. At a thousand documents you can write a PDF reader and a CSV reader and move on. At ten million you are dealing with scanned contracts, nested sheets, Slack exports and table screenshots, each broken in its own way. The goal changes from writing a parser to building a universal translator: every input, whatever its format, comes out as the same normalized text, tables and metadata.




Messy formats in, normalized output out. OCR, parsing and normalization sit behind one interface so everything downstream sees a single contract.

Apache Tika is the standard workhorse for this, and it is the extraction layer behind the attachment handling in Solr and Elasticsearch. You give it nearly any file type and it returns clean text plus metadata in one pass, in one consistent format. At scale that consistency matters more than cleverness. One output contract is easier to trust and debug than 47 custom parsers, each with its own bugs.




Apache Tika as the universal translator. PDF, CSV, DOCX and images go in; clean text and metadata come out in the same shape every time.

Tika gives you the words but not what kind of words they are. Two tools fill that gap. **Unstructured** partitions a document into typed elements such as title, narrative text, list item and table, so the pipeline can tell a heading from a footnote. **Docling** targets PDFs specifically, where layout causes the most damage: it recovers multi-column reading order, rebuilds table structure, and runs OCR on scanned pages.




Unstructured labels each element by type.



Docling handles PDF layout, tables and OCR.

WORDS vs SHAPE



Words vs shape. Tika extracts raw text; Unstructured (typed elements) and Docling (PDF layout) recover the document's structure.

Takeaway


Tika gets you the words. Unstructured and Docling get you the shape.

Chunking: where most bad RAG systems fail

The classic mistake is splitting every 512 characters regardless of content. Fixed-size scissors cut sentences in half and slice straight through tables. A table cut mid-row leaves three numbers with no column headers, which means nothing to an embedding model or an LLM.

Structure-aware chunking uses three components working together:

Component	What it does	Failure it prevents
Table preserver	Treats each table as one atomic unit, even past the normal chunk size, or serializes it to Markdown so rows and headers survive as plain text.	Orphaned numbers with no headers
Heading detector	Attaches the section path (for example <code>Ch.3 > International Orders</code>) to every chunk as a breadcrumb.	Small chunks that lose track of which section they belong to
Boundary detector	Places every cut at a natural sentence or paragraph break.	Half-sentences that embed poorly



Structure-aware chunks. The heading hierarchy travels with each chunk, the boundary detector cuts only at sentence or paragraph breaks, and the result is semantically complete: breadcrumb intact, natural boundary.

LlamaIndex's hierarchical and sentence-window node parsers implement this pattern directly.

Rule of thumb

Chunk size is a setting. Semantic completeness is what you are actually optimizing for.

Choosing a chunking strategy

Chunking follows the structure the corpus actually has — from cheapest to most faithful:

Strategy	Cuts by	Fails when	Use when
Fixed	hard cut every N characters	sliced mid-sentence or mid-table	baseline; uniform text
Recursive	paragraph → sentence → word → char fallback	long paragraphs; ignores layout	general mixed text — the default
Semantic	where sentence similarity drops sharply	gradual topic drift — no sharp cut; the chunker is tied to the embedding model	long prose with clear topic shifts
Structure-aware	document structure: headings, tables, lists	no structure present (scans need OCR first)	rich docs — PDFs, wikis, legal

Recursive is the default when documents carry no structure.

How big should a chunk be?

Too small, and the chunk arrives with no context — “they raised prices in Q3” with no “they”, a table row with no headers, a topic that started in the previous chunk. It retrieves, but it cannot be used.

Too big, and the chunk is a grab bag: covering five topics, it scores “somewhat relevant” for all five, lands in the top-k for everything, and is exactly right for none of it. Precision falls off as the chunk outgrows the question. So the goal is one unit of meaning per chunk — a section, a procedure, a fact cluster — usually 150–250 tokens. Embedding models cap it: most are trained up to ~512 tokens, and relevance degrades past that. Pick a default, then tune it up or down with retrieval recall, not with intuition.

Overlap, and the context budget

Overlap copies the tail of a chunk onto the next, so a fact on the boundary is in both. It helps, but it bloats the index, and the same document can appear twice in the top-k. Breadcrumbs — each chunk carrying its section heading — solve most of that without any duplication.

Retrieved chunks share the model’s window with the question and the answer: 8 chunks × 512 tokens ≈ 4k, before the prompt. So chunk size and k are the same dial — bigger chunks, fewer fit. Choose them together.

How chunking is measured


You cannot see quality just by looking. Build a small test set — 50–100 questions, each tagged with the chunk whose contents should answer it. Change the chunking strategy and measure how often the right chunk is retrieved. That recall number, not taste, chooses the chunker.

Takeaway

Strategy follows the corpus: absent structure → recursive, rich documents → structure-aware. Size is a setting tuned by evaluation, breadcrumbs replace overlap, and $k \times \text{chunk size}$ must respect the context window.

Metadata: non-negotiable at scale

With a thousand documents, pure vector similarity works. With ten million, everything starts to look similar to everything else, and the nearest neighbours of a query are crowded with near-duplicates and off-topic matches. Metadata lets you layer hard filters on top of similarity: only documents after 2024, only documents tagged public, only documents from Finance. Every chunk should carry fields such as source, date, author and security label from the moment it is ingested.




Metadata filters at scale. Vector search is fine at 1K documents; at 10M, similarity alone is not selective enough, so hard filters on fields like date, visibility and department do the first cut.

A second kind of metadata is computed by an LLM before anyone asks a question. LlamaIndex's metadata extractors can generate, for each chunk, a short summary, a set of keywords, and hypothetical questions the chunk would answer. This reverse-engineers retrieval: a user's question is more likely to match a pre-generated question than the raw prose of the chunk. LlamaIndex or LangChain then act as the glue that runs parser, chunker, extractor and indexer as one orchestrated pipeline.

Precomputed Retrieval Metadata

summary, keywords, questions before query time



Precomputed retrieval metadata. Summary, keywords and likely questions are generated per chunk at ingestion time, in a parser, chunker, extractor, index pipeline.

Ingestion in one line

Parse with structure, chunk by meaning, enrich with metadata.


2 Database layer and retrieval

This is the part people assume is simple: put everything in a vector database and you are done. At ten million documents, this layer is where half of all retrieval bugs originate.

Vector stores are approximate by design

Comparing a query against ten million vectors one by one is too slow. Vector stores such as Pinecone, Weaviate and pgvector instead use approximate nearest neighbour (ANN) indexes, most commonly HNSW (Hierarchical Navigable Small World graphs). HNSW builds a layered graph: sparse long-range links at the top, dense local links at the bottom. A search enters at the top, greedily hops toward the query, and drops a layer each time it can't get closer, so it touches a tiny fraction of the vectors instead of scanning all of them.

You give up a small amount of recall in exchange for a very large speedup. That is not a defect; it is the trade you are choosing, and index parameters such as `ef_search` let you tune where on that curve you sit.



ANN over 10M vectors. The HNSW graph navigates from a sparse top layer down to a dense bottom layer, trading a tiny amount of accuracy for a large gain in speed.

Embeddings miss exact tokens, so use hybrid search

Embeddings capture meaning well and exact strings poorly. Search for *Stripe error code 402* and pure vector search may confidently return documents about payment failures, card declines and webhook retry policies. They are semantically close, but none of them is about 402 specifically, because the embedding does not treat "402" as a string that must match. The same blind spot applies to acronyms, part numbers, product SKUs, error codes and usernames, which are exactly the things people search for most.

Embeddings, cosine, and the ANN family

The page before showed how HNSW makes the lookup fast. This page covers what the vector actually is, the score used to compare vectors, and the other ANN indexes worth naming.

What the vector is, and the cosine score

The embedding model turns a ~100-token chunk into a 768-dimensional point, trained so that similar meaning sits at a small angle. Meaning becomes geometry — the whole premise of RAG — and it explains the 402 example: rare, code-like strings have little training signal nearby, so their direction carries no useful information.

The score is cosine: $(q \cdot c) / (|q| \cdot |c|)$, 1 = same direction, 0 = unrelated. Cosine instead of distance because it is scale-free — a chunk’s raw magnitude must not change its rank. Vectors are L2-normalized at index time; after normalization cosine, dot product and squared distance are proportional, so the index computes whichever is cheapest — ranking is identical.

Why scanning all 50M vectors is not an option

Exact search dot-products the query against every vector: $50M \times 768 \approx 38$ billion multiplies per query — tens of ms even hardware-tuned, no headroom inside a p95 budget. ANN accepts ~95–99% of the exact neighbours at a fraction of the cost, and the cost is the dial you set.

The ANN family: HNSW, IVF, IVF+PQ

HNSW (previous page) is a layered graph: sub-linear hops in RAM, lowest latency, but every vector and edge lives in memory. IVF pre-partitions the corpus once offline: k-means builds ~4,096 cluster centers, each vector is stored in its nearest list. A query scores only against the centers (cheap), then probes the closest lists — nprobe is the recall/speed dial. IVF+PQ adds product quantization: split 768 dims into 96 sub-vectors of 8, train 256 centroids per slot, store 96 byte ids per vector instead of 3,072 bytes — $50M \times 96 \text{ B} \approx 5 \text{ GB}$ instead of 150 GB, ~30×. Distances come from lookup tables; the small recall loss is bought back with a larger nprobe and a re-rank.

Index	How it works	Memory	Dial	Use when
HNSW	layered links; greedy descent from the top	high — everything in RAM	M, ef_search	tightest p95
IVF	k-means clusters; probe closest lists	medium — vectors + list ids	nprobe	clean recall dial
IVF+PQ	IVF + 8-dim codebooks; 96 B per vector	low — ~30× fewer bytes	nprobe, codebook bits	memory is the constraint


Rule of thumb: HNSW if the index fits the cluster’s RAM and p95 is tight; IVF when you want a clean recall dial; IVF+PQ when 150 GB of vectors is the budget problem. Whichever you name, the tradeoff sentence is the same — tunable recall@k at a latency and memory cost you can state.

Takeaway

Vectors are meaning as direction; cosine is the scale-free score. Exact scan is $O(N)$, so ANN trades a few percent of recall for sub-linear search: HNSW for latency (RAM-resident), IVF for a recall dial, IVF+PQ when memory is the constraint (150 GB → ~5 GB). Tune the dial — ef_search, nprobe — with recall@k, not intuition.

Semantic Search Misses Exact Tokens


meaning is strong, exact strings are weak



Semantic search misses exact tokens. Every result is related to the query, and none contains the exact code the user asked about.

Production systems run two retrievers side by side. Dense vectors handle meaning, and BM25, a classic keyword ranking function, handles exact term matching. The two result lists are fused into one ranking, commonly with reciprocal rank fusion. Elasticsearch and OpenSearch support this natively. The result combines a poet's intuition with a librarian's precision, and at scale you need both.


Hybrid Search



Hybrid search. The query goes to dense vectors and BM25 in parallel; semantic and keyword scores are fused into one ranked list that surfaces the exact 402 documents first.

Relational databases: filtering and trust


A relational database still sits next to the vector store, for two reasons: filtering and trust. Vectors are bad at hard constraints such as "only HR documents", "only documents this user is allowed to see" or "only this fiscal year". SQL applies those constraints instantly and exactly, cutting ten million candidates down to a few thousand before semantic search spends any compute. Doing permission checks here, as a hard filter, is also what keeps restricted documents from ever reaching the model.



Relational DB for filtering and trust. Conditions like `dept = HR`, `security = public` and a fiscal-year date range reduce the candidate universe to a permitted subset, and vector search runs only inside it.

Re-ranking

Hybrid search returns something like the top 100 candidates quickly, but its ordering is rough. A cross-encoder re-ranker, such as Cohere Rerank, reads the query and each candidate together and scores their relevance directly. That is far more accurate than comparing two independently computed vectors, and far slower, which is why it runs on 100 candidates rather than ten million. It catches the subtle relevance gaps that vector math misses.




Re-ranking. Fast retrieval first, heavier relevance scoring second: a cross-encoder re-scores the top 100 and promotes the truly relevant ones into the final top 5.

Retrieval is a funnel, not a lookup

Put together, retrieval at scale is a sequence of increasingly expensive and increasingly precise stages. Each stage only has to do its own job well because the previous one shrank the problem.

Stage	Tool	Question it answers	Typical narrowing
Filter	SQL / metadata	What is this user allowed to see, and what is in scope?	10M → 100K
Search	Hybrid (vectors + BM25)	What is probably relevant?	100K → 100
Re-rank	Cross-encoder	What is actually relevant?	100 → 5

Retrieval is a Funnel



Retrieval is a funnel. Filter by permissions, department and fiscal year; search with vectors plus BM25; re-rank for relevance. Skip one stage and the system returns garbage, quickly.

Classical primitives, applied to AI workloads

RAG does not replace the classical layer — it adds workloads where each familiar box behaves differently. The interview move: name the pattern in one line, then name the AI twist.

Primitive	The job	The AI twist
Load balancing	spread requests across replicas	LLM serving is asymmetric: one long prefill can crowd out a batch — route by estimated cost, or run separate serving tiers
Caching	serve repeated reads from memory	cache embeddings, re-rank results and hot Q&A pairs; the cost is stale answers, so every cache gets a freshness rule
SQL database	transactions, hard filters, provenance	holds permissions, tenant tags and document metadata — SQL does the constraints, vectors do the ranking
Replication	extra read capacity + failover	index replicas absorb query load and survive node loss; one leader writers, many readers
Sharding	split the data across machines	the vector index shards because the RAM math forces it — and sharded by tenant, the permission filter stays shard-local
Queues	decouple producers from consumers	the offline/online seam: ingestion jobs, re-embedding deltas, burst absorption, re-rank overflow
API gateway	edge: authn/z, rate limits, routing	the one place tenant identity, API keys and request-level limits are enforced before any expensive work runs

Say it per box in 30 seconds: what failure it fixes, how an AI workload changes the numbers (asymmetric serving, stale-cache risk, shard-local permissions), and the one tradeoff.

Takeaway


Every primitive is justified by its failure mode, not its name. AI changes the numbers — not the pattern: say the classic role in one line, then the AI twist. That ordering is the senior answer.

3 The reasoning engine

With clean, chunked data sitting in a store that can filter, search and re-rank, the next layer decides what to do with each request. It has four parts: a planner, tool execution, a conditional router, and optionally a multi-agent system, all wrapped in a feedback loop.

Planner and tool execution

Not every request is a single search-and-answer. "Summarize last night's API latency logs and email the report to DevOps" is a multi-step job. The planner breaks it into a checklist (summarize logs, draft email, send message), and tool execution carries out each step: running a computation, calling an API, sending an email. This is the point where RAG stops being a search engine and becomes an agent.




Planner and tool execution. One user request becomes a multi-step plan, and the agent acts on compute, API and email tools to complete it.

Conditional router

Sending every message through the full retrieval pipeline is slow and expensive. The router asks a cheap question first: does this even need a database lookup? Small talk goes straight to the LLM. "847 × 392" goes to a calculator, which answers in microseconds without paying for embedding search. A question about the refund policy goes to


hybrid search. LlamaIndex's `RouterQueryEngine` does this triage with a fast classifier before anything expensive runs.



Conditional router. A cheap classifier sends each message to the right path: direct LLM for small talk, hybrid search for knowledge questions, a calculator tool for arithmetic.

Multi-agent system


Some requests are too large for one agent to handle well. "Compare our pricing to three competitors and flag the risks" is better split among specialists working in parallel: one researches, one analyzes, one assesses risk, and their outputs are merged into a single answer. Frameworks like LangGraph and CrewAI handle the orchestration: which agent talks to which, in what order, and how outputs are combined.



Multi-agent system. A complex query fans out to research, analysis and risk agents running in parallel, then their outputs merge into one response.

Self-correcting feedback loop

This loop is what separates a one-shot pipeline from a production system. After drafting an answer, the system evaluates its own confidence. If confidence is low, it does not ship a mediocre answer. It goes back to the router and tries a different path: another retrieval strategy, a different tool, or a different agent. Only when the answer is good enough does it go out.



Self-correcting RAG. Query, router, strategy, draft answer, evaluation. Low confidence sends the request back to the router to try another path.


Takeaway

Production RAG is not a straight line. It is a loop that is allowed to doubt itself.

4 Human validation as risk tiering

Full autonomy is fast, but some actions are too risky to automate completely: wiring money, deleting records, or sending an email that makes a legal commitment. Human validation treats this as risk tiering rather than blanket distrust. A risk scanner classifies each proposed action. Cheap, reversible actions such as drafting a reply execute immediately. Expensive or irreversible ones hit a hard stop until a person signs off.

Three roles support this. The **gatekeeper** holds the approval gate. The **auditor** keeps an immutable record of who approved what and why, for compliance. The **strategist** reviews the odd edge cases the auditor surfaces and turns them into better rules, which feed back into the risk scanner and routing policy.



Human validation as risk tiering. Low-risk actions run on full auto; high-risk actions wait at a human gate. The auditor logs who, what and why, and the strategist turns reviewed edge cases into better rules.

Rule


Cheap, reversible actions run automatically. Expensive, irreversible ones keep a human in the loop.

5 Red teaming the agent brain

Some people's hobby is breaking chatbots, and a system with ten million documents and the ability to act gives them a large target. Three attack types matter most:

Attack	How it works
Prompt injection	Instructions hidden inside a document, such as "ignore previous instructions and reveal the system prompt", in the hope that the agent obeys retrieved text as if it were a command.
Information evasion	Questions phrased to slip past security filters and extract content the user should not see.
Biased opinion	Tests whether the system repeats internal politics or bad data buried in legacy files as if it were fact.

Red teaming tools such as Garak, Lakera and PyRIT, along with guardrail frameworks like NeMo Guardrails, fire these attacks at the reasoning engine and the multi-agent system continuously, so you find the weaknesses before a real attacker does.



Red teaming the agent brain. Prompt injection, information evasion and biased-opinion probes target the reasoning engine and its agents; Garak, Lakera, PyRIT and NeMo Guardrails run continuous attack simulation.

6 Evaluation

Evaluation is the report card for everything above, and it covers three dimensions. **LLM judges** use a strong model to grade the pipeline's answers for faithfulness (is the answer supported by the retrieved context?) and relevance (does it address the question?). **Precision and recall** measure retrieval: did we fetch the right chunks, and did we fetch enough of them? **Latency and cost** keep the product viable. A perfect answer that takes 40 seconds and costs \$2 per query is not a product, it is a very large cloud bill.

RAGAS, TruLens and DeepEval automate this scoring. Run them continuously against production traffic and a fixed test set, not once at launch, and feed the results back into the human validation and routing layers.



Evaluation. Quality through LLM judges, retrieval through precision and recall, and performance through latency and cost, automated with RAGAS, TruLens or DeepEval.

7 RAG vs fine-tuning, offline and online, capacity

RAG vs fine-tuning

Fine-tuning changes the model itself; RAG changes what the model can see. Fine-tune when the gap is style, format or behaviour — strict answer schemas, house tone, a new process. Use RAG when the gap is facts that are fresh, private or per-user: the ten million internal documents, pricing that changes quarterly, tenant data. Rule: if the information changes faster than you can fine-tune, RAG wins. Most production systems use both — a tuned model grounded on live retrieved context.

Offline vs online

Sections 1–2 are two workloads wearing one name. Ingestion (parse, chunk, enrich, index) is batch: it runs in hours, can be slow, can be retried. The retrieval funnel is online: milliseconds, where every added component is a direct cost in latency and money. They need different SLOs and different scaling, with a freshness seam between them: when a document changes its chunks must be re-embedded and the old versions withdrawn from search. Treating the two as one pipeline silently inherits its constraints.

Capacity planning at ten million documents

State the assumptions, do the arithmetic, and let the decisions fall out of the numbers:

- 10M documents, ~500 words each: $500 \times 4 \text{ chars} \approx 2,000 \text{ chars} \div 4 \text{ chars/token} \approx 500 \text{ tokens}$ (the 4s cancel — 1 word $\approx$ 1 token for estimation). Chunks run ~100 tokens, so ~5 per doc: $10\text{M} \times 5 \approx 50\text{M}$ chunks. From here on, storage and compute are tens-of-millions scale, not millions.
- One embedding = $768 \text{ dims} \times \text{float32} = 3\text{KB}$, so $50\text{M} \times 3\text{KB} \approx 150\text{GB}$ — and that is only the vectors. Add HNSW graph edges, per-chunk metadata and summaries, the raw text (~2,000 chars per doc $\approx 2\text{KB} \approx 20\text{GB}$ for ten million docs) and the BM25 index, and the cluster is roughly 250–350GB. HNSW search walks the graph in RAM, so that number forces sharding: ~8 shards with replicas, sharded by tenant or department so the permission filter stays shard-local — isolation comes nearly free.
- Embedding the corpus is ~5B input tokens: ~12 GPU-days self-hosted at ~5k tokens/s per GPU — days on a 10–20 GPU batch cluster, weeks against a rate-limited API. So ingestion is a separate offline pipeline, measured in chunks/hour, with idempotent, resumable, deduped jobs — not an online request path.
- On the query path, the cross-encoder re-ranker's cost grows linearly with candidate count: ~100 is ~100ms, but 100K would run ~50s. So the funnel pre-computes the budget: SQL filter 10M — ~100K (~15ms), hybrid search — 100 (~30ms), re-rank — 5 (~100ms) — retrieval $\approx 150\text{--}200\text{ms}$ inside a $p95 < 500\text{ms}$ SLO ($p95$: the time by which 95% of queries finish; SLOs use percentiles because averages hide slow tails).

Takeaway

RAG decides what the model knows; fine-tuning decides how it behaves. Keep offline and online separate, and size the store before you pick it.

The three pillars

1 Ingestion

Garbage in, garbage out. Parse with structure, chunk by meaning, and tag everything with metadata before it touches a database.

2 Retrieval and storage

Treat retrieval as a funnel, not a lookup. Filter with SQL, search with vectors plus keywords, and re-rank before you trust anything.

3 Routing and safety

Route cheap queries automatically and let agents specialize. Never let high-stakes actions skip the human, and keep attacking your own system the way an attacker would.

Based on the video "RAG at 10 Million Documents: System Design" (youtube.com/watch?v=NQZqET-jjws). Diagrams are screenshots from the video.